

Understanding the Dangers of Concurrency

Joe Hummel, PhD

joe-hummel@pluralsight.com

<http://blog.joehummel.net/>

[@joehummel](#)



Overview



- **Your presenter: Joe Hummel, PhD**
 - PhD in field of high-performance computing
 - An exciting time to be working in this area...

- **Agenda for this module:**
 - *The many pitfalls when parallel programming*
 - *The most common pitfall — race conditions*
 - *Locking vs. Lock-free*
 - *Synchronization primitives*
 - *Concurrent Data Structures*
 - *Lots of demos*

***Beware* the pitfalls of concurrency**

- Async and Parallel programming are full of potential pitfalls...
 - *Race conditions*
 - *Starvation*
 - *Livelock*
 - *Deadlock*
 - *Optimizing compilers*
 - *Optimizing hardware*
 - For a quick read on the subject:
 - "Tools And Techniques to Identify Concurrency Issues", R. Patil and B. George, MSDN Magazine, June 2008
- Online: <http://msdn.microsoft.com/en-us/magazine/cc546569.aspx>

Correctness?

- Your code should provide 2 guarantees:

Safety: *nothing bad happens.*

Liveness: *eventually something good happens.*

Important Terminology

Def: a race condition exists when the outcome depends on the timing of events.

*Race conditions
are bad !*

Def: a critical section is smallest region of code involved in a race condition.

*Identify and
resolve...*



The **danger** of shared resources

- Most common **pitfall** — concurrent access to **shared resources**

- variables
- objects
- collections
- files
- ...

```
int sum = 0; // Parallel code:

Task.Factory.StartNew(() =>
    { sum = sum + obj1.Computation(); });

Task.Factory.StartNew(() =>
    { sum = sum + obj2.Computation(); });

Task.Factory.StartNew(() =>
    { sum = sum + obj3.Computation(); });
```

Danger == concurrent

Safe == concurrent



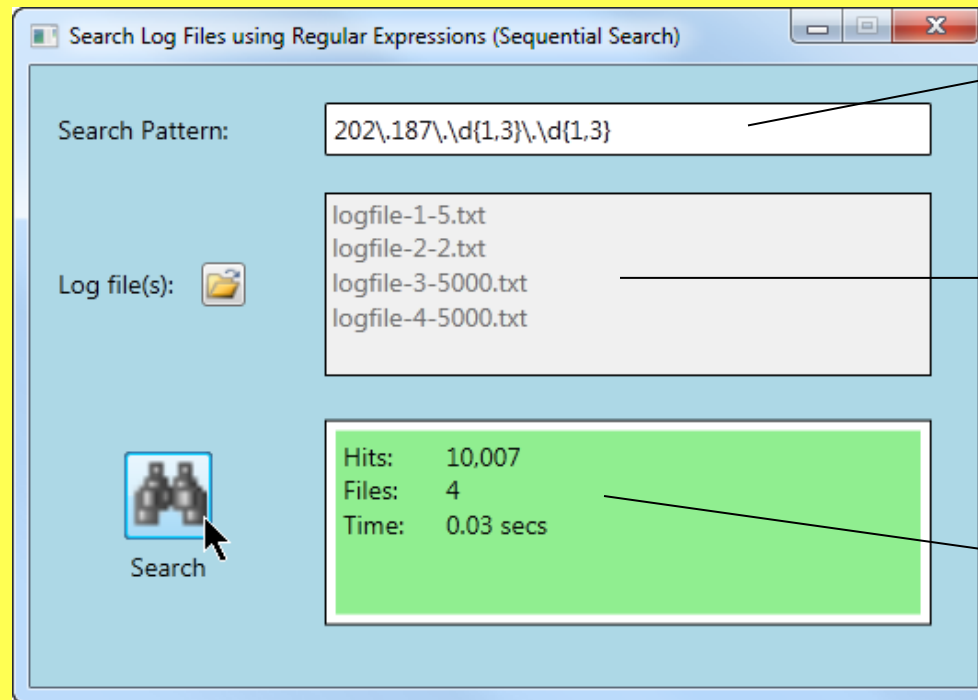
Error!!

Danger == concurrent reads & writes, or concurrent writes;

Safe == concurrent reads.

DEMO

- Searching log files...



search string as a regular expression (e.g. IP address 202.187.*.*)

files to search...

results

Solutions?

- **Redesign** to eliminate shared resources / critical sections
 - e.g., each task uses only local memory (aka "*task local storage*")
- Use **thread-safe** entities within critical sections
 - e.g., TPL offers thread-safe data structures (*System.Collections.Concurrent*)

ConcurrentBag

ConcurrentDictionary

BlockingCollection

ConcurrentStack

ConcurrentQueue

- Use **synchronization** to control access within critical sections

CountdownEvent / ManualResetEvent / AutoResetEvent

Lock

Monitor

Semaphore

Barrier

Interlocked

Mutex

SpinLock / SpinWait

Solution #1: Lock

- Surround critical section with a **lock** on a common object

- restricts entry to one task at a time*

Parallel

Sequential

```
int sum = 0;

sum = sum + obj1.Computation();
sum = sum + obj2.Computation();
```

```
int sum = 0;
var l = new object(); // common object:


Task.Factory.StartNew(() =>
{
    int t = obj1.Computation();
    lock(l) { sum += t; }
});

Task.Factory.StartNew(() =>
{
    int t = obj2.Computation();
    lock(l) { sum += t; }
});
```

Locking is a general technique based on .NET Monitor class

DEMO

- Searching log files in **parallel** using a **lock**...



```
lock (1)
{
    hits++;
}
```

Solution #2: Interlocked

- A HW-based lock for simple, arithmetic critical sections

- *typically more efficient*

Sequential

```
int sum = 0;

sum = sum + obj1.Computation();
sum = sum + obj2.Computation();
```

Parallel

```
using System.Threading;
```

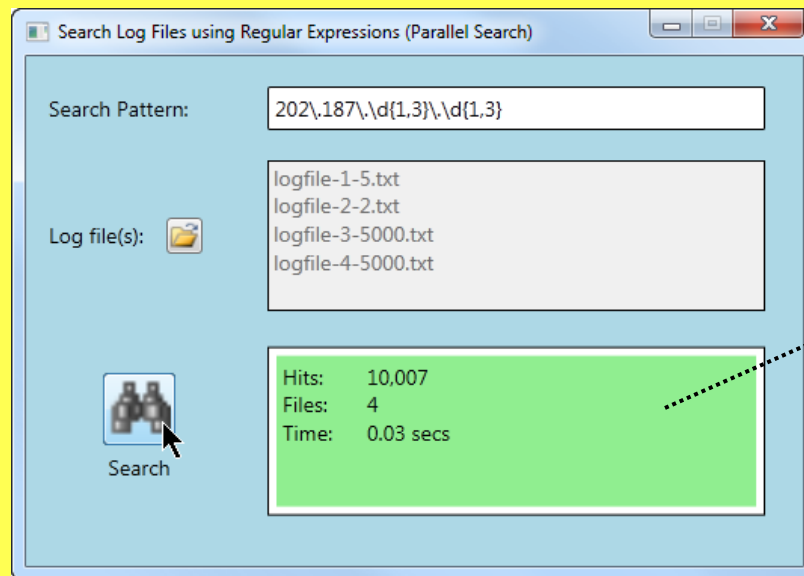
```
int sum = 0;
```

```
Task.Factory.StartNew(() =>
{
    int t = obj1.Computation();
    Interlocked.Add(ref sum, t);
});
```

```
Task.Factory.StartNew(() =>
{
    int t = obj2.Computation();
    Interlocked.Add(ref sum, t);
});
```

DEMO

- Searching log files in **parallel** using **interlocking**...



`Interlocked.Increment(ref hits);`

Solution #3: Lock-free

- Locking leads to **contention**, hindering performance
- Use **lock-free**¹ designs whenever possible...
- Example:
 - use **local resource** instead of shared resource
 - **return** result instead of updating shared resource
 - join, **combine** partial results into final result
 - no locking in this case

Local memory ("task local state")

```
var t1 = Task.Factory.StartNew<int>(() =>
{
    int t = obj1.Computation();
    return t;
});

var t2 = Task.Factory.StartNew<int>(() =>
{
    int t = obj2.Computation();
    return t;
});

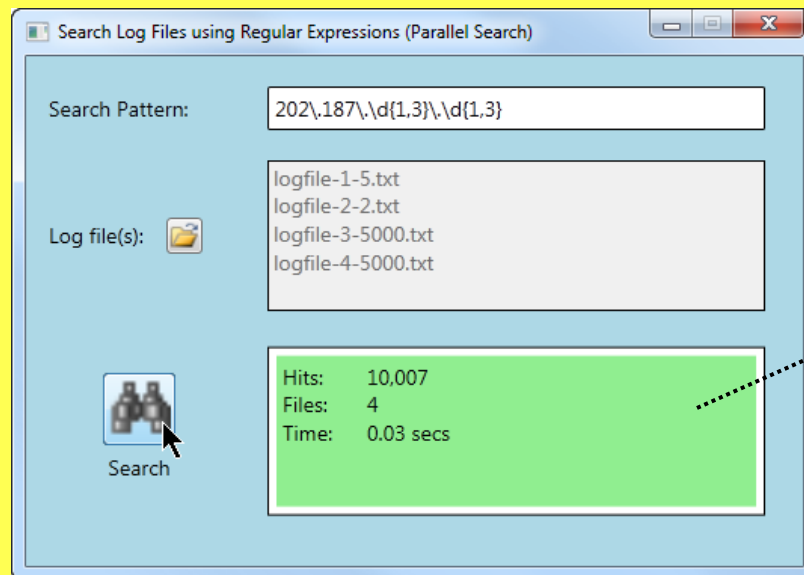
sum = t1.Result + t2.Result;
```

Wait... Harvest partial results,
produce final result

¹Lock-free is not necessarily freedom from locks, but freedom from arbitrary waiting...

DEMO

- Searching log files in **parallel** using a **lock-free** approach...



```
int lHits = 0;  
.  
.  
.  
lHits++;  
.  
.  
.  
return lHits;
```

Another danger...

- Another form of race condition:

- *shared objects and **non-thread-safe** classes...*

Random and List are
not thread-safe...

Sequential

```
List<int> results = new List<int>();
Random rand = new Random();

for (int i=0; i < N; i++)
{
    int r = RunSimulation( rand );
    results.Add(r);
}
```

Parallel

```
List<int> results = new List<int>();
Random rand = new Random();

for (int i=0; i < N; i++)
{
    Task.Factory.StartNew(() =>
    {
        int r = RunSimulation( rand );
        results.Add(r);
    })
}
```

Def: an object or method is thread-safe if parallel use does not cause a race condition.

Error!!



Solution?

- *Locking is one solution...*
- *Other solutions?*

Sequential

```
List<int> results = new List<int>();
Random rand = new Random();

for (int i=0; i < N; i++)
{
    int r = RunSimulation(rand);
    results.Add(r);
}
```

thread-safe data structure

```
using System.Collections.Concurrent;

var results = new ConcurrentQueue<int>();

for (int i=0; i < N; i++)
{
    Task.Factory.StartNew(() =>
    {
        // use random seed so tasks generate different sequences:
        var rng = new RNGCryptoServiceProvider();
        byte[] data = new byte[4];
        rng.GetBytes(data);
        int seed = BitConverter.ToInt32(data, 0);
        var rand = new Random(seed);

        var rand = new Random();

        int r = RunSimulation(rand);
        results.Enqueue(r);
    });
}
```

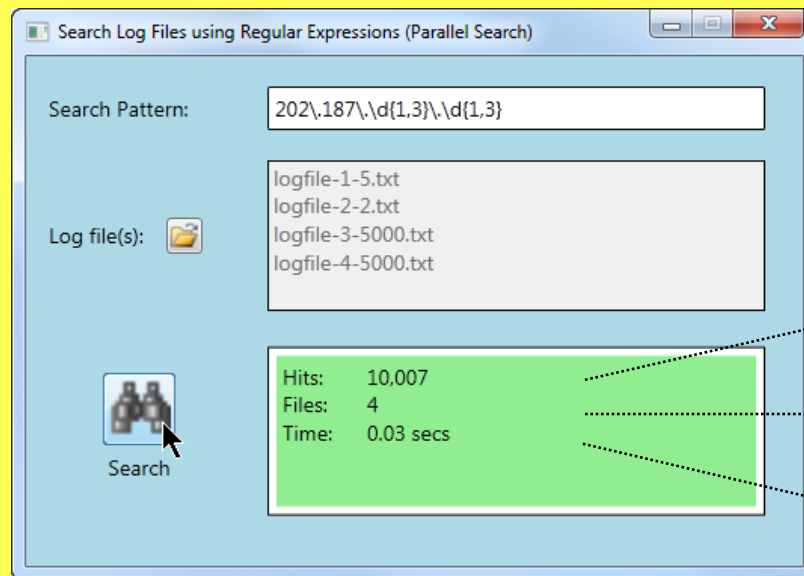
yields duplicates

local instance

Correct solutions can be subtle, and hard to test!

DEMO

- Thread-safety when searching log files in parallel...



single, shared Regex object? ✓

File.ReadAllBytes? ✓

Encoding.UTF8.GetString? ?

Synchronization primitives in brief...

- TPL builds upon existing .NET synchronization primitives:


Primitive	Purpose
Monitor	<i>General-purpose .NET synchronization class</i>
Lock	<i>Enforces one-at-a-time semantics using Monitor class</i>
Mutex	<i>Win32 lock suitable for inter-process sync ("mutual exclusion")</i>
Interlocked	<i>Hardware-based lock for simple, arithmetic operations</i>
Semaphore	<i>Enforces N-at-a-time semantics via Win32</i>
SpinLock / SpinWait	<i>Lock-like mechanisms that loop ("spin") instead of yield CPU</i>
Barrier	<i>Allows tasks to synchronize ("sync-up") before start of next phase</i>
CountdownEvent, ManualResetEvent, AutoResetEvent	<i>Additional ways for tasks to synchronize with one another...</i>

*Goal of TPL is to provide better abstractions
so you *don't* need all these...*



Concurrent Data Structures

- TPL offers a set of thread-safe data structures:
 - `ConcurrentBag<T>`
 - `ConcurrentDictionary<T>`
 - `ConcurrentQueue<T>`
 - `BlockingCollection<T>`
 - `ConcurrentStack<T>`



*Redesigned and optimized for
concurrent access --- we'll discuss in
future lecture ("Designs and Patterns")*

A few words about performance...

- Future lecture on *Designs and Patterns* will discuss in more detail
- But there's a few lessons now worth mentioning...

I/O is **hard** to parallelize

- *Requires support from OS and HW*
- *Don't expect much from commodity HW*

Experiment...

- *Try different ideas*

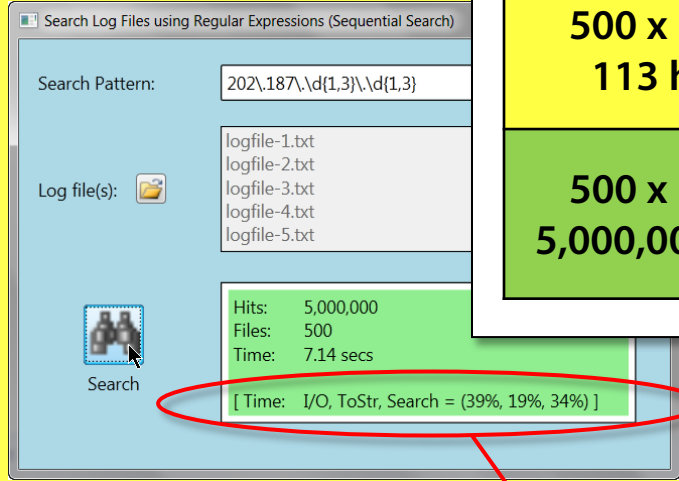
Instrument your app

- *Time things*
- *Measure things*

Amdahl's law

- *Performance limited by sequential component*

DEMO: performance of log file search app



Input	Sequential	Parallel-LockFree	
		<i>version 1</i>	<i>version 2</i>
500 x 1MB 113 hits			
500 x 1MB 5,000,000 hits			

shared RE object

local RE object

Instrumentation: % of I/O, ToString, Search

Summary


- Beware the many dangers of async and parallel programming:

RACE CONDITIONS **LIVELOCK** **DEADLOCK**
OPTIMIZING COMPILERS
STARVATION **OPTIMIZING HARDWARE**

- The most common pitfall: **race conditions from shared resources**

- Various solutions...

- *redesign to eliminate*
- *use thread-safe entities*
- *use synchronization*



References

- **Microsoft's main site for all things parallel:**
 - <http://msdn.microsoft.com/concurrency>
- **MSDN technical documentation:**
 - <http://tinyurl.com/pp-on-msdn>
- **Background reading on the dangers of concurrency, and solutions:**
 - Just about any textbook on *Operating Systems*
 - *Concurrent Programming on Windows*, by Joe Duffy, Addison-Wesley
 - *Win32 Multithreaded Programming*, by A. Cohen and M. Woodring