

Execution Model and Types of Parallelism

Joe Hummel, PhD

joe-hummel@pluralsight.com

<http://blog.joehummel.net/>

[@joehummel](#)



Overview



- **Your presenter: Joe Hummel, PhD**
 - PhD in field of high-performance computing
 - An exciting time to be working in this area...

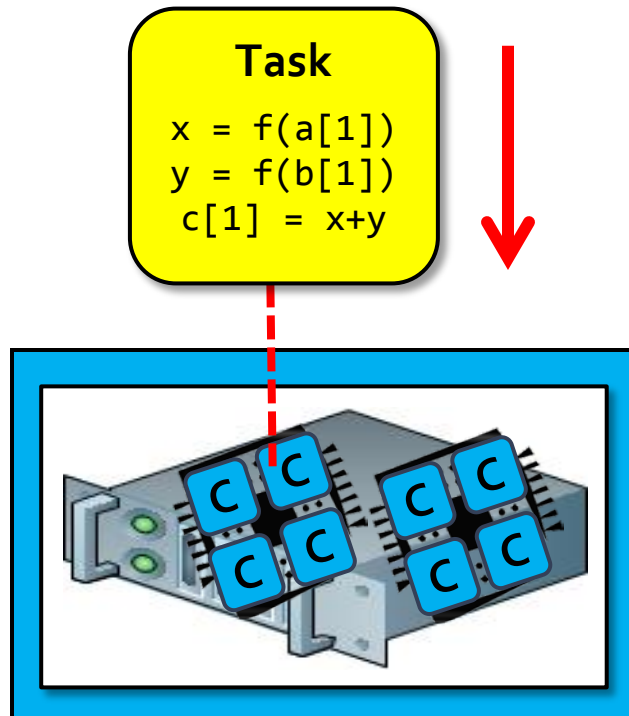
- **Agenda for this module:**
 - *Execution model*
 - *Tasks vs. Threads vs. Cores*
 - *Types of parallelism:*
 - *data, task, dataflow, embarrassingly parallel*
 - *TPL support:*
 - *Parallel.For, .ForEach, .Invoke, and more*

Tasks, tasks, tasks

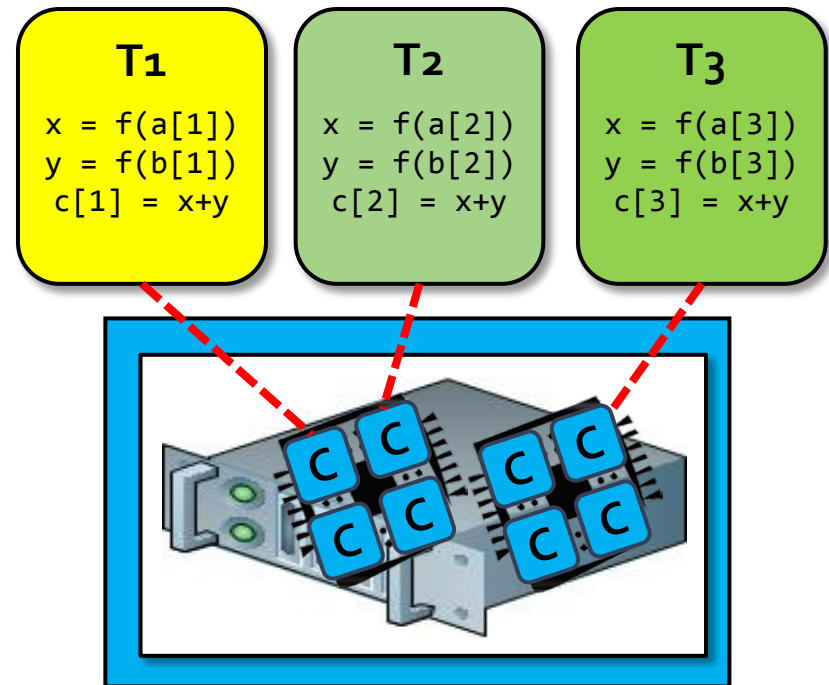
- Our job as developers is to express work as tasks
- .NET's job is to execute these tasks intelligently
- Let's drill-down on the task execution model in .NET 4...

Execution model

- A task executes sequentially:



- Multiple tasks may execute in parallel:



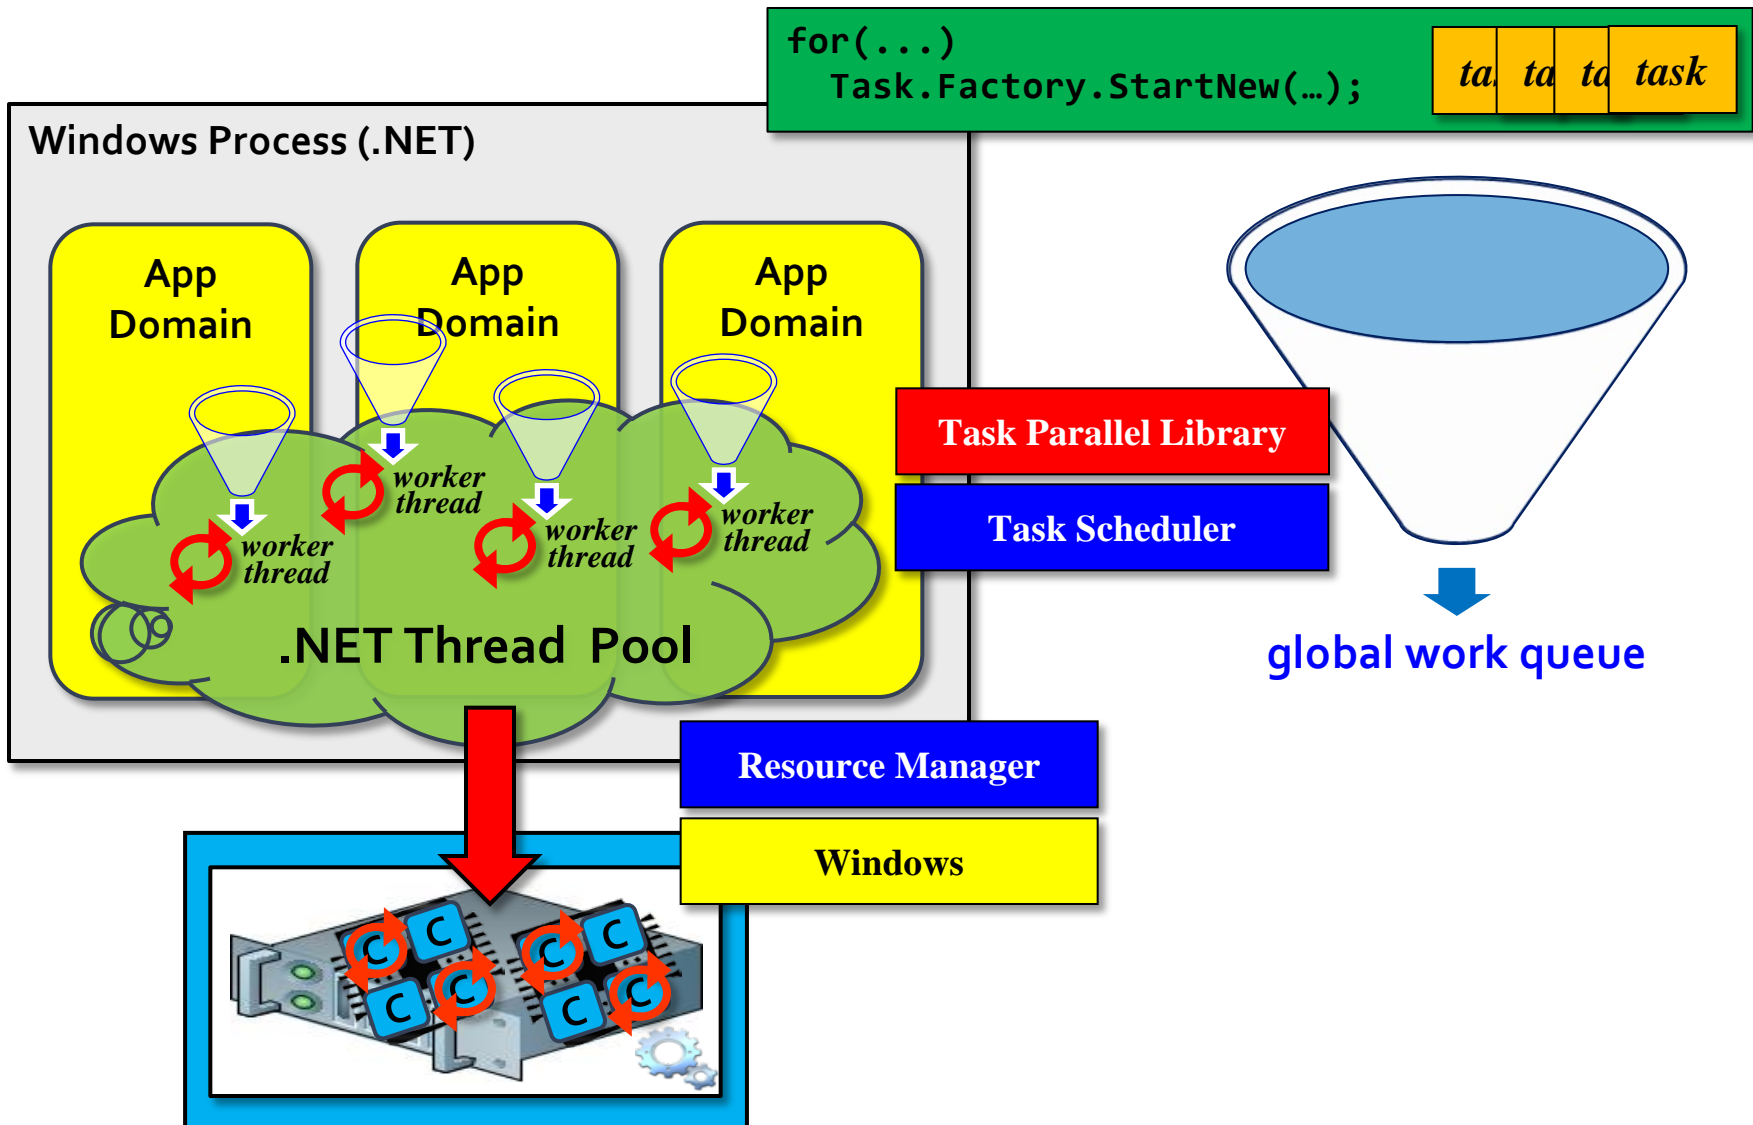
Task granularity

- Tasks are lightweight objects
 - *but there's still a cost associated with creating & manipulating*
- What is **minimum granularity** to offset this cost?

```
Task t = Task.Factory.StartNew( () =>  
    {  
        :  
        :  
        :  
    }  
);
```

*Task execution should take
at least 200-300 hundred
CPU cycles...*

Tasks vs. Threads vs. Cores



Want to customize task scheduler?

- Add task priorities?
- Use threads from your own thread pool?
- Supply your own task scheduler...

```
TaskScheduler myTS = new MyTaskScheduler();  
TaskFactory   myTF = new TaskFactory(myTS);  
  
Task t = myTF.StartNew( () => {...} );
```

```
class MyTaskScheduler : TaskScheduler  
{  
    .  
    .  
    .  
}
```

For more details and examples, see:

<http://msdn.microsoft.com/en-us/library/ee789351.aspx>

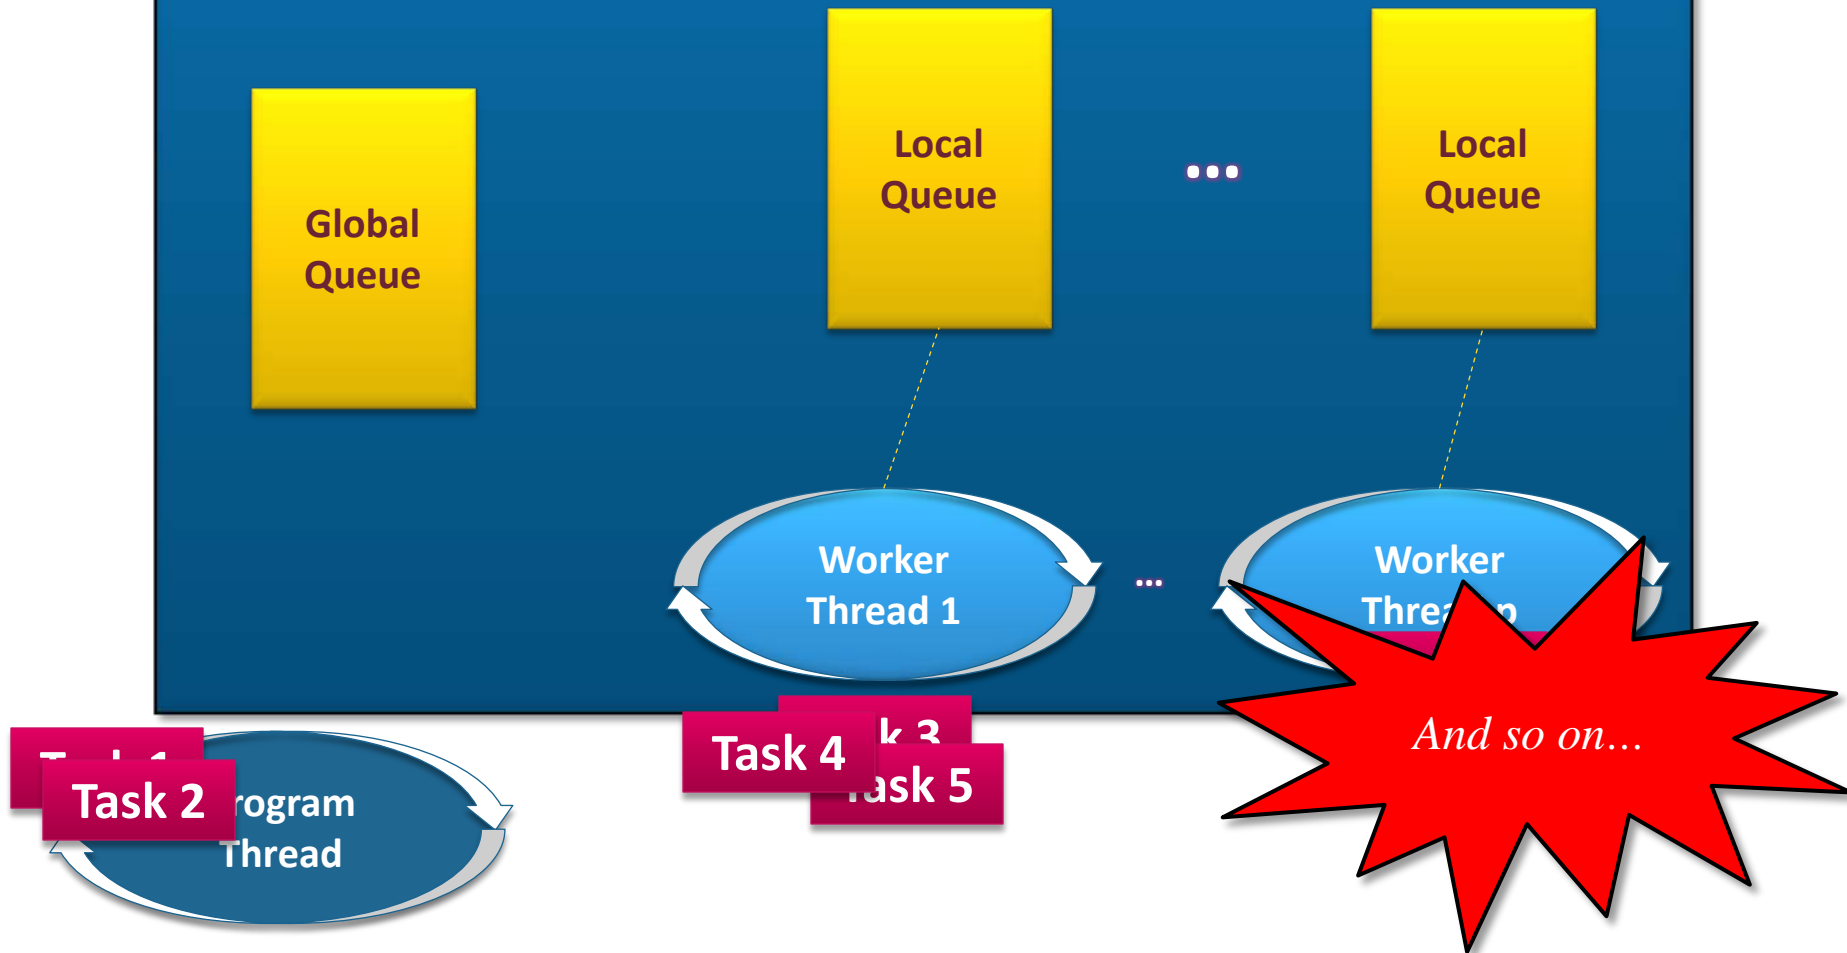
<http://code.msdn.microsoft.com/ParExtSamples>

Default task scheduler — observations...

- A single, global work queue does not scale
 - ==> *local queues reduce contention*
- Responds dynamically to demand
 - *a thread is dedicated to a particular tasks until task completes – fast execution*
 - *if task code blocks & waits, its thread blocks & waits – not available to others*
 - ==> *creates new worker threads to maintain task completion rate*
- Dynamically balances the workload
 - ==> *"work-stealing" of tasks from other queues for load-balancing*

Work-stealing

Thread Pool



Task scheduler assumptions

(1) *Tasks are short-lived — at most 1-2 seconds.*

(2) *Execution order doesn't matter — random order of execution is okay.*

What if execution order matters?

- If you want to execute tasks in order of creation...
- ... create tasks with **fairness** option:

```
Task.Factory.StartNew(  
    () => { /* long running code */ },  
    TaskCreationOptions.Fairness  
);
```

.NET adds task to global queue, never to worker's local queue...

Guarantees that tasks will **start** in order they are created, but tasks could finish in **any** order

Slight increase in cost — global queue carries higher synchronization overhead.

What if task is long-running?

- Where "long" ==> greater than a few seconds...
- ... create task with long-running option:

```
Task.Factory.StartNew(  
    () => { /* long running code */ },  
    TaskCreationOptions.LongRunning  
);
```

.NET creates a non-worker pool thread, dedicates thread to this task...

Beware of cost — roughly 200,000 cycles to create, 100,000 to retire, and every context switch is 6,000-8,000. Plus memory cost for stack space.

DEMO

- Suppose we have 100 long-running computations to perform...

What happens if we start 100 tasks with no creation options?

Every couple seconds a new thread is created and a task started...



What happens if we start 100 tasks with long-running option?

Very quickly, 100 threads are created and 100 tasks started...

If computations are CPU-intensive, system becomes over-subscribed with too many threads competing for cores...

Solution?

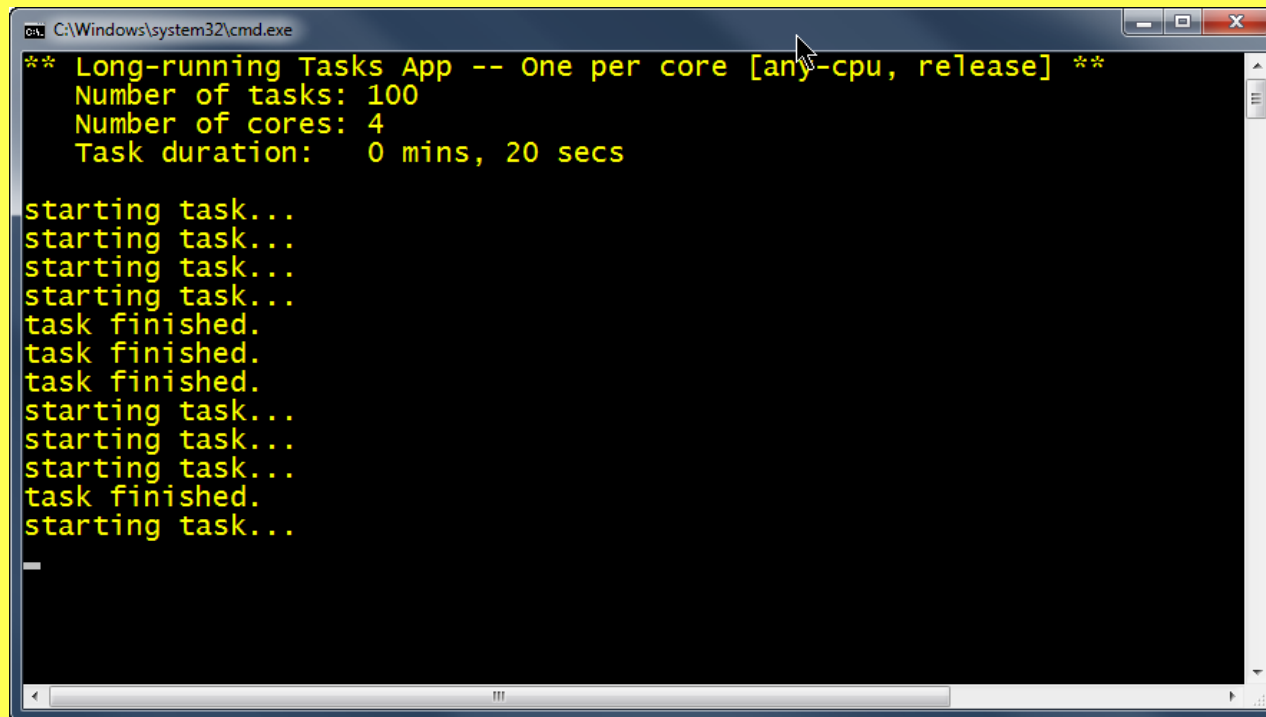
- It's okay to create hundreds or thousands of tasks...
- It's **not** okay if those tasks are long-running, CPU-intensive tasks
 - *you have too many threads competing for cores*
- **Solution?**
 - *Start just T tasks — one per core*
 - *As tasks finish, start new ones...*

```
int numCores = System.Environment.ProcessorCount;
for (int i=0; i<numCores; ++i)
    tasks.Add( Task.Factory.StartNew( () => {...} ));

while (tasks.Count > 0) // wait for tasks to finish:
{
    int index = Task.WaitAny( tasks.ToArray() );
    tasks.RemoveAt(index);
    if (still work to do)
        tasks.Add( Task.Factory.StartNew( () => {...} ));
}
```

DEMO

- Correctly executing 100 long-running computations...



```
C:\Windows\system32\cmd.exe
** Long-running Tasks App -- one per core [any-cpu, release] **
Number of tasks: 100
Number of cores: 4
Task duration: 0 mins, 20 secs

starting task...
starting task...
starting task...
starting task...
task finished.
task finished.
task finished.
starting task...
starting task...
starting task...
task finished.
starting task...
```

Aside — here's an even better way...

- Use **parallel for loop** with option to constrain amount of parallelism
 - *A teaser for what's coming next :-)*


```
using System.Threading.Tasks;  
  
var options = new ParallelOptions();  
options.MaxDegreeOfParallelism = System.Environment.ProcessorCount;  
  
Parallel.For(0, N, options, (i) => { /*long-running computation*/ });
```


Common types of parallelism...


- Data
- Task
- Dataflow
- Embarrassingly parallel

(1) Data parallelism

- **Def:** the same operation executed across different data.



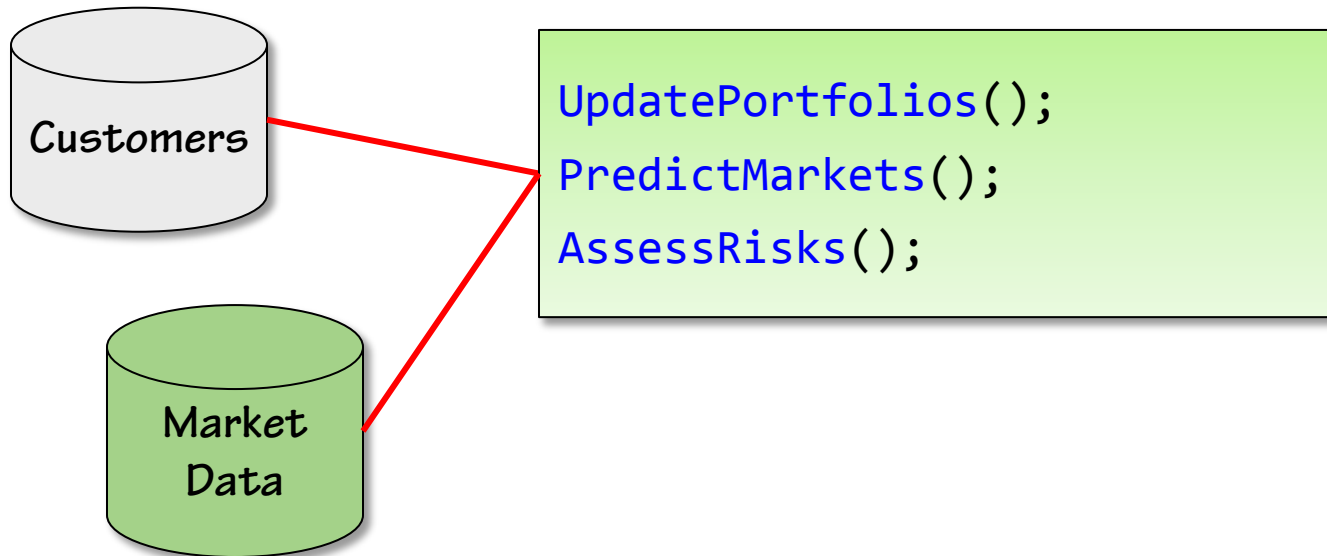
```
foreach(Customer c in Customers)  
  UpdatePortfolio(c);
```



```
for(i=0; i<N; i++)  
  for(j=0; j<N; j++)  
    A[i,j] = sqrt(c * A[i,j]);
```

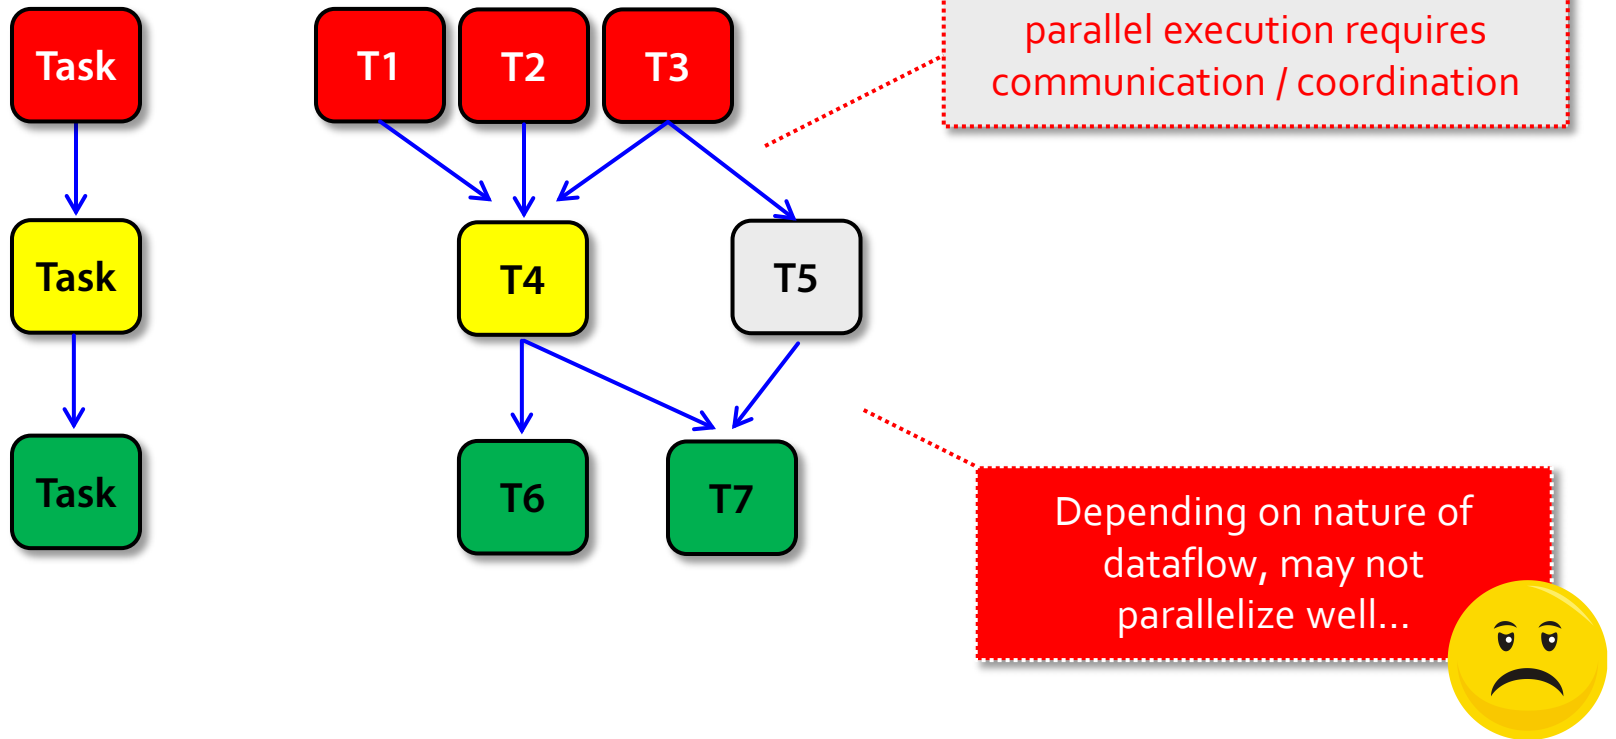
(2) Task parallelism

- **Def:** different operations executed across the same or different data.



(3) Dataflow parallelism

- **Def:** when operations depend on one another.
 - *data "flows" from one operation to another*




(4) Embarrassingly Parallel

- **Def:** a problem is embarrassingly parallel if the computations are entirely independent of one another.

*==> *no* data flow*

Original code

```
foreach(Customer c in Customers)  
    UpdatePortfolio(c);
```



```
Parallel.Foreach(Customers, (c) =>  
{  
    UpdatePortfolio(c);  
});
```

*Also referred to as
"delightfully parallel", best
type of parallelism!*




TPL support

- **Dataflow:**
 - *next lecture "Designs and Patterns"*
- **Data parallelism:**
 - `Parallel.For`
 - `Parallel.Foreach`
- **Task parallelism:**
 - `Parallel.Invoke`


Parallel class

```
for(i=0; i<N; i++)  
    DoWork(i);
```



```
Parallel.For(0, N,  
    (i) =>  
    {  
        DoWork(i);  
    }  
);
```

```
foreach(var e in ds)  
    DoWork(e);
```



```
Parallel.ForEach(ds,  
    (e) =>  
    {  
        DoWork(e);  
    }  
);
```

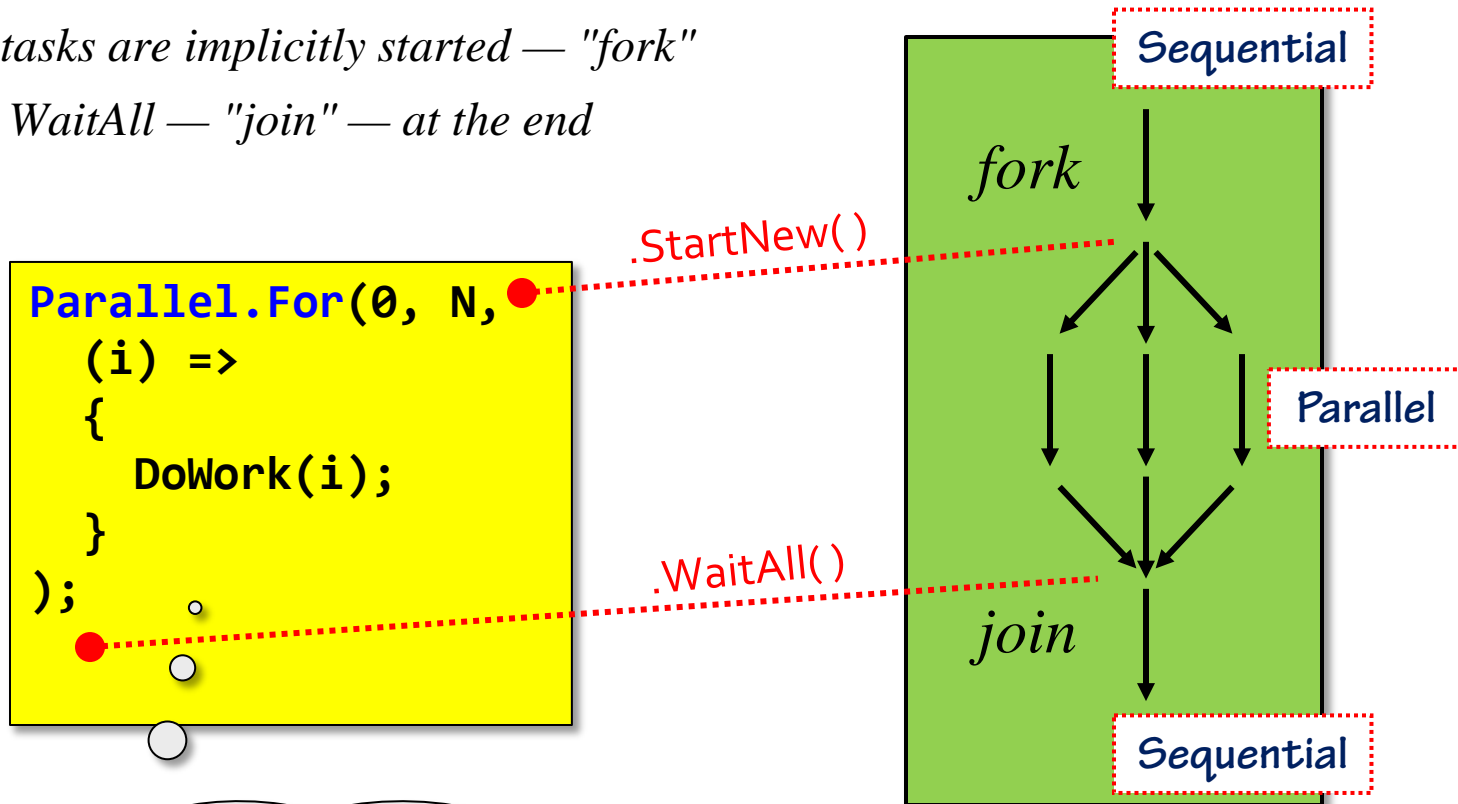
```
Task1();  
Task2();  
Task3();
```



```
Parallel.Invoke(  
    () => Task1(),  
    () => Task2(),  
    () => Task3()  
);
```

Structured parallelism

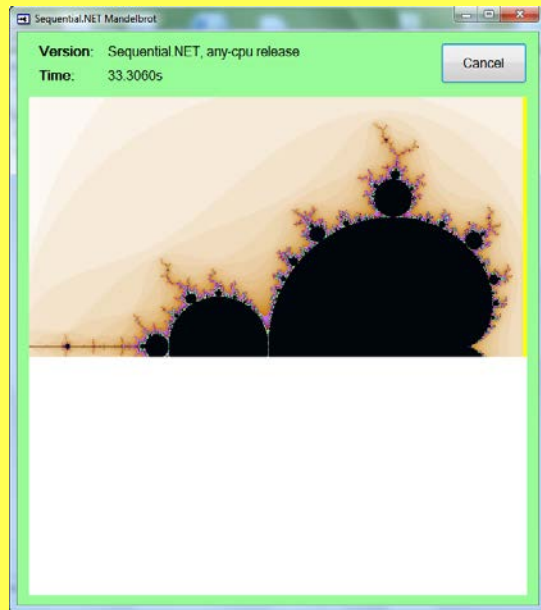
- **Parallel** class uses "*fork-join*" pattern
 - *a set of tasks are implicitly started* — "fork"
 - *implicit WaitAll* — "join" — *at the end*



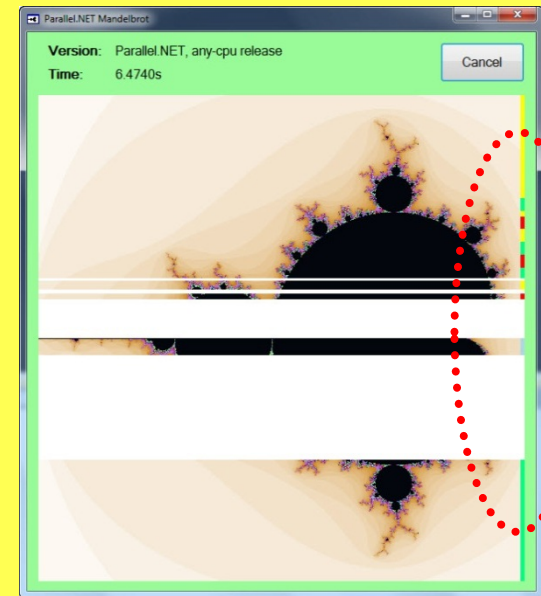
When in doubt, use this pattern —
easiest to understand & apply...

DEMO


- Exploiting data parallelism in Mandelbrot app



Parallel.For



Data partitioning




- **Parallel.For & .Foreach "chunk" the data for good load-balancing**
 - *chunks start small, then grow in size*
 - *assumes loop body does significant computation to offset cost of small chunks*
 - *not so good for simple loop bodies --- use something else?*

Custom data partitioning

- Parallel.Foreach supports **custom partitioners**
- Example:
 - *partitioning simpler loop bodies for efficient parallelization*

```
for (i=0; i<N; ++i)  
{ /*simple*/ }
```



```
using System.Collections.Concurrent;  
using System.Threading.Tasks;  
  
Parallel.Foreach( Partitioner.Create(0, N), /*range-like*/  
    (range) =>  
    {  
        for (int i=range.Item1; i<range.Item2; ++i)  
        { /*simple*/ }  
    });
```

For more details and examples, see:

<http://msdn.microsoft.com/en-us/library/dd997411.aspx>

<http://www.drdoobs.com/go-parallel/article/224600406>

<http://code.msdn.microsoft.com/ParExtSamples>

Exception handling

- How are exceptions handled by the **Parallel** class?
- If an unhandled exception is thrown by `.For` / `.Foreach` / `.Invoke`, then
 - tasks *allowed to finish* their current iteration / work
 - exception(s) are *deferred and re-thrown* as `AggregateException`

```
try
{
    Parallel.For(0, N, (i) =>
    {
        .
        .
        .
    })
}
catch(AggregateException ae)
{ /*handle*/ }
```


throw new ApplicationException("oops");

Break out of loop — from ***within***

- You can break out of Parallel.For / .Foreach
- Task(s) finish current iteration, then exit
 - call **Stop** if you want to stop loop ASAP
 - call **Break** if you want earlier iterations to complete (0..M-1)

```
ParallelLoopResult lr = Parallel.For(0, N, (i, loopControl) =>
{
    if (some condition becomes true)
        loopControl.Stop();
    else
        DoWork(i);
});

if (!lr.IsCompleted)
    // loop did not run to completion:
```



Cancel loop — from ***outside***

- Task(s) finish current iteration, then exit

- beware: an exception is thrown!*

```
void DoLoop(CancellationTokenSource cts)
{
    var options = new ParallelOptions();
    options.CancellationToken = cts.Token;

    try {
        Parallel.For(0, N, options, (i) =>
        {
            DoWork(i);
        });
    }
    catch(OperationCanceledException oce) { /*loop canceled*/ }
    catch(AggregateException ae)          { /*something else*/ }
}
```

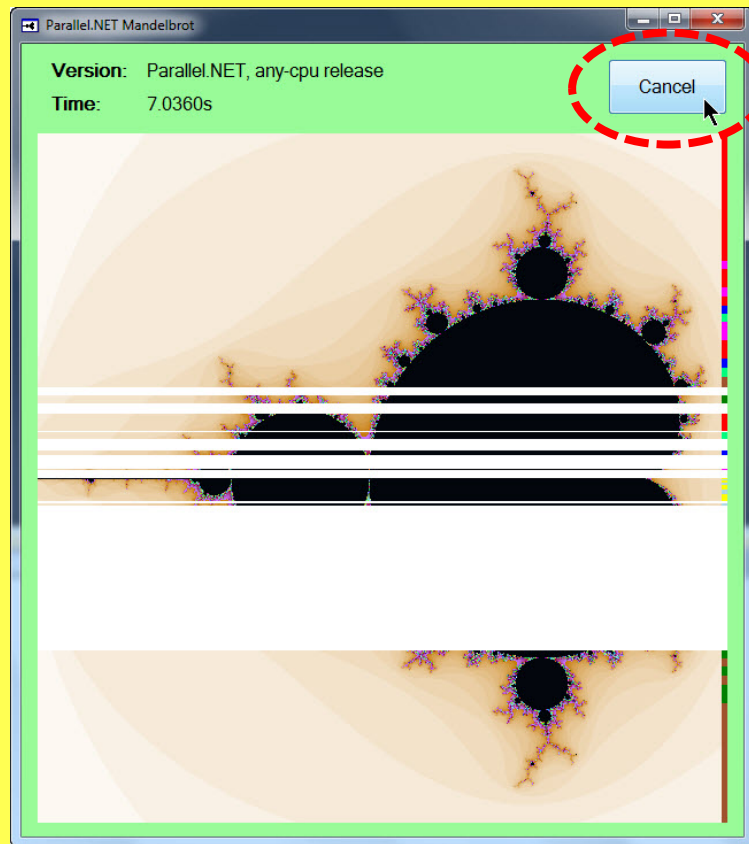
```
void StartButton_Click(...)
{
    var cts =
        new CancellationTokenSource();

    Task.Factory.StartNew(
        () => { DoLoop(cts); });
}

void CancelButton_Click(...)
{
    cts.Cancel(); // stop loop:
}
```

DEMO

- Mandelbrot app with **cancellation**...



Summary

- **Your job** is to create tasks
- **.NET's job** is to execute those tasks efficiently
- **Task Parallel Library** provides a versatile, efficient execution engine
 - *load-balancing task scheduler*
 - *intelligent resource manager*
- **Identify the type of parallelism in your application:**
 - Data, Task, Dataflow, Embarassingly Parallel
- **Use higher-level abstractions where appropriate:**
 - `Parallel.For`, `.Foreach`, `.Invoke`

References

- **Microsoft's main site for all things parallel:**
 - <http://msdn.microsoft.com/concurrency>
- **MSDN technical documentation:**
 - <http://tinyurl.com/pp-on-msdn>
- **I highly recommend the following short, easy-to-read book:**
 - *Parallel Programming with Microsoft .NET: Design Patterns for Decomposition and Coordination on Multicore Architectures*, by C. Campbell, R. Johnson, A. Miller and S. Toub, Microsoft Press

Online: <http://tinyurl.com/tpl-book>