

# Designs and Patterns for Parallel Programming

---

Joe Hummel, PhD

[joe-hummel@pluralsight.com](mailto:joe-hummel@pluralsight.com)

<http://blog.joehummel.net/>

[@joehummel](#)



# Overview



- **Your presenter: Joe Hummel, PhD**
  - PhD in field of high-performance computing
  - An exciting time to be working in this area...
  
- **Agenda for this module:**
  - *A few high-level design problems...*
  - *Pipeline and Dataflow*
  - *Using the Concurrent Data Structures provided by the TPL*
  - *Producer-Consumer*
  - *MapReduce and Task Local State*
  - *Parallel LINQ*
  - *Speculative Execution*
  - *Asynchronous Programming Model*

# Design Problem #1

*Your app needs to perform 100+ CPU-intensive operations, each taking roughly 3-5 minutes.  
Execution order doesn't matter.*

*How do you execute the 100+ operations?*



Create 100 tasks, 1 per op,  
with no creation options.

*[ Since long-running, .NET will over-compensate and create 100 worker threads; since CPU intensive, system will thrash. ]*

Create 100 tasks, 1 per op,  
with long-running option.

*[ .NET creates 100 dedicated threads; since CPU intensive, system will thrash. ]*

Create 1 task per core, as one  
finishes create another.

*[ use WaitAllOneByOne pattern, or  
MaxDegreeOfParallelism. ]*

# Design Problem #2

*Your app needs to perform 20+ I/O operations — e.g. web page downloads. Order doesn't matter, but downloads can take anywhere from a few ms to a few seconds.*

*How do you download the web pages?*



✓ Create 1 code task per download, no special creation options.

*[ Starts one download for each worker thread in pool. ]*

✓ Create 1 façade task per download using FromAsync + APM pattern (web object's Begin/End methods).

*[ Starts download and returns thread to pool, potentially allowing all 20+ downloads to start. ]*

✓ Create 1 task per core, explicitly or via MaxDegreeOfParallelism.

*[ Starts one download per core. ]*


✓ Use Producer-Consumer pattern, e.g. dedicating 1-2 tasks to download & remainder to processing.

*[ Consider when download time  $\ll$  processing time. ]*

# Design Problem #3

*You want a logging task that runs for the duration of your app, logging say every 30 seconds.*

*What does the design look like?*



**This is harder than it looks :-)**

*Things to consider in your design:*

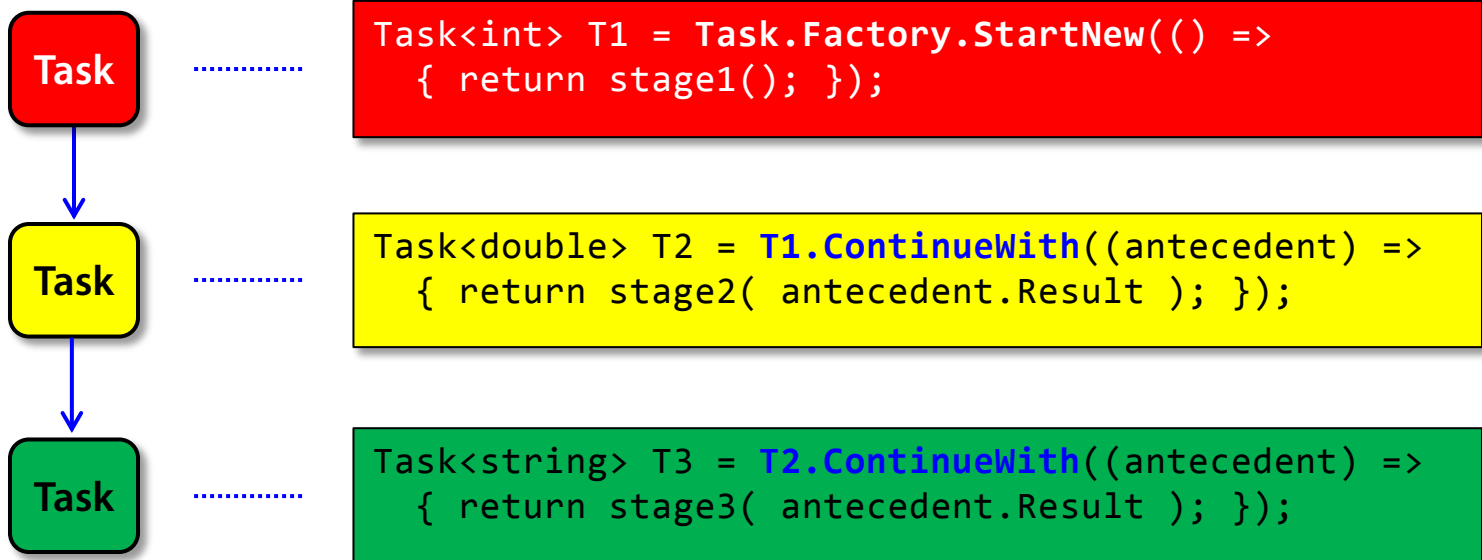
1. *Create with long-running option*
2. *Before app closes you need to join with task and catch any exceptions*
3. *Design clean shutdown of task so app closes (use TPL cancellation?)*
4. *If task can crash, design a way to monitor and restart (check task's **Status** property via a timer, or app's message loop?)*

# Patterns

- Pipeline
- Dataflow
- Concurrent data structures
- Producer-Consumer
- Map-reduce
- Parallel LINQ
- Speculative Execution
- Asynchronous Programming Model


# Pipeline

- **Linear flow** from one task to another
  - *image processing, UI updating, workflows...*



# Dataflow

- Generalized flow with many-to-one and one-to-many...



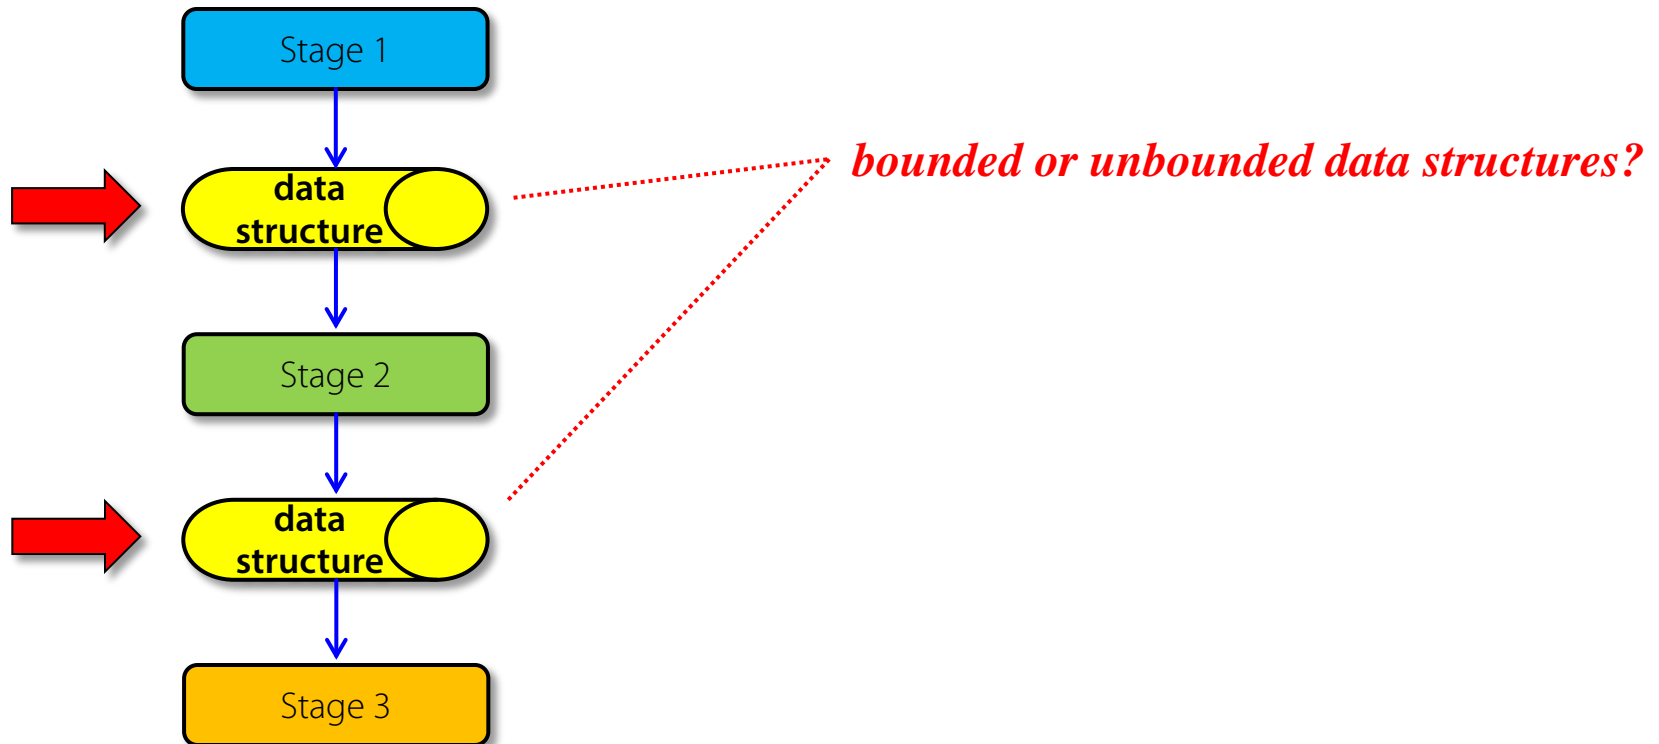
```
Task<int> T1 = ...;
Task<int> T2 = ...;
Task<int> T3 = ...;

Task<double> T4 =
    Task.Factory.ContinueWhenAll<int, double>(
        new[] {T1, T2, T3},
        (tasks) =>
        {
            double result;
            foreach(Task<int> t in tasks)
                ...
            return result;
        }
    );
```



# Increasing parallelism

- Thread-safe data structures allow stages to run **independently**...



# Concurrent Data Structures

- TPL offers a set of thread-safe data structures:

- `ConcurrentBag<T>`
- `ConcurrentQueue<T>`
- `ConcurrentStack<T>`
- `ConcurrentDictionary<T>`

- `BlockingCollection<T>`

Bounded in size, blocking tasks  
when full or empty



*Designed and optimized for  
concurrent access...*

# Example: ConcurrentQueue

## ■ Enqueue?

- *safe to call in parallel...*

## ■ Dequeue?

- *typically you check then dequeue...*
- *this is \*still\* unsafe: you cannot separate check from dequeue, some other task could grab in-between — yielding race condition!*
- *Solution? API redesigned with **TryDequeue** method that checks and dequeues in one call*

## ■ Enumeration?

- *Foreach, ToArray, ...*
- *thread-safe by creating snapshots*

```
using System.Collections.Concurrent;

var Q = new ConcurrentQueue<int>();

Task.Factory.StartNew(() =>
{
    Q.Enqueue(123);
});
Task.Factory.StartNew(() =>
{
    Q.Enqueue(456);
});

Task.Factory.StartNew(() =>
{
    if (Q.Count > 0)
        DoWork(Q.Dequeue());
});

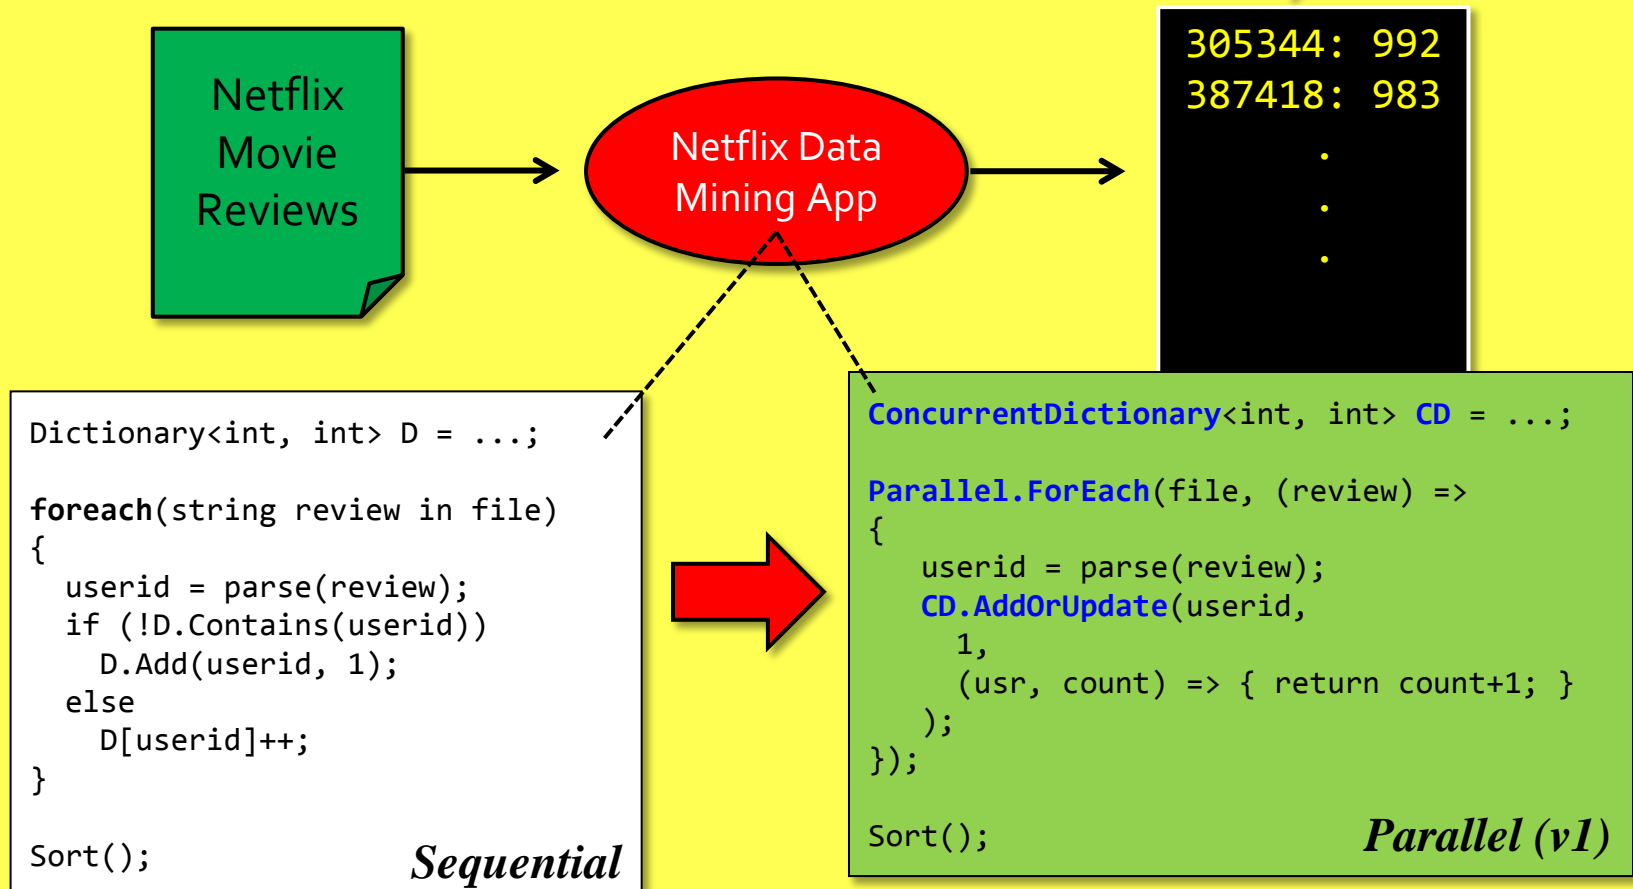
Task.Factory.StartNew(() =>
{
    Work w;
    if (Q.TryDequeue(out w))
        DoWork(w);
});

Task.Factory.StartNew(() =>
{
    contents = Q.ToArray();
});
```

# DEMO

*Users with most reviews  
(the "top 10" users)*

- Netflix data mining using a **ConcurrentDictionary**...



# Results?

- Netflix app worked **correctly** with ConcurrentDictionary




- Unfortunately, parallel version ran **slower**



- *In this application, minimal data processing ==> contention for dictionary*
- *ConcurrentDictionary does internal synchronization / locking*
- *Too much contention ==> slower execution...*

# Producer-Consumer

- Good pattern to use for long-running workloads where **data generation** speed is very different from **data consumption** speed
  - *Data structure **throttles** faster component...*
- Example:
  - **producer(s)** read from disk / network
  - **consumer(s)** process the data




# BlockingCollection<>

- Think **fixed-sized collection** in a **parallel** world:

- *blocking producers if collection is full*
- *blocking consumers if collection is empty*

- Fixed-sized collections are more realistic


- ... and typically *\*improve\* performance by throttling faster component!*



# Implementation

```
int maxCapacity = ...; // collection is a fixed-size:
var workQ       = new BlockingCollection<T>(maxCapacity);

// Create a new factory to produce long-running tasks:
var tf = new TaskFactory( TaskCreationOptions.LongRunning,
                          TaskContinuationOptions.None );
```



```
Task producer = tf.StartNew(() =>
{
    for(...) // blocks task if Q is full:
        workQ.Add( work );

    // Signal we're done:
    workQ.CompleteAdding();
});
```

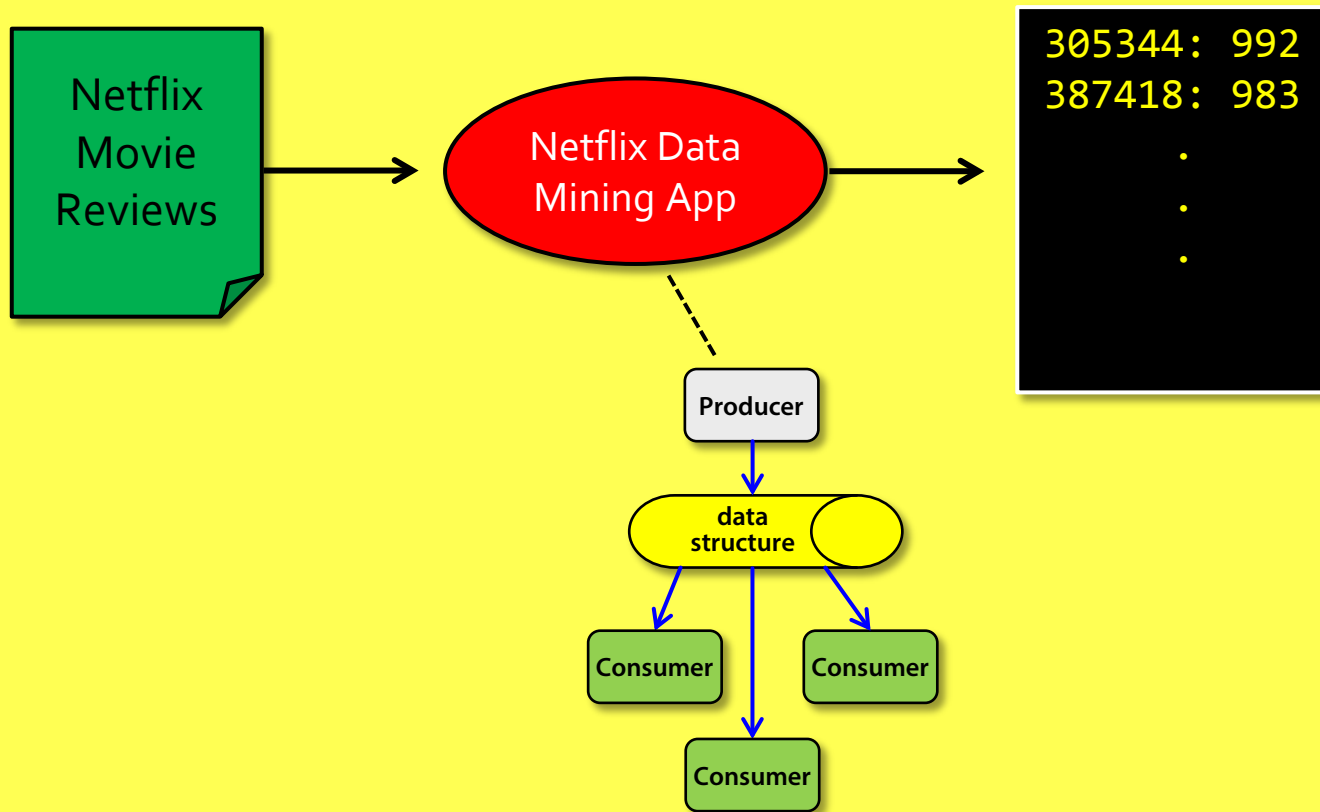
```
Task consumer = tf.StartNew(() =>
{
    while (!workQ.IsCompleted)
    {
        try { // blocks task if Q is empty:
            T work = workQ.Take();
            : // process work item:
            :
        }

        // Exceptions ==> no more work:
        catch(ObjectDisposedException) { }
        catch(InvalidOperationException) { }
    }
});
```



# DEMO (v2)

- Netflix data mining using **Producer-Consumer**...



# Results?

- Producer-Consumer was correct \*and\* yielded a performance **improvement!**



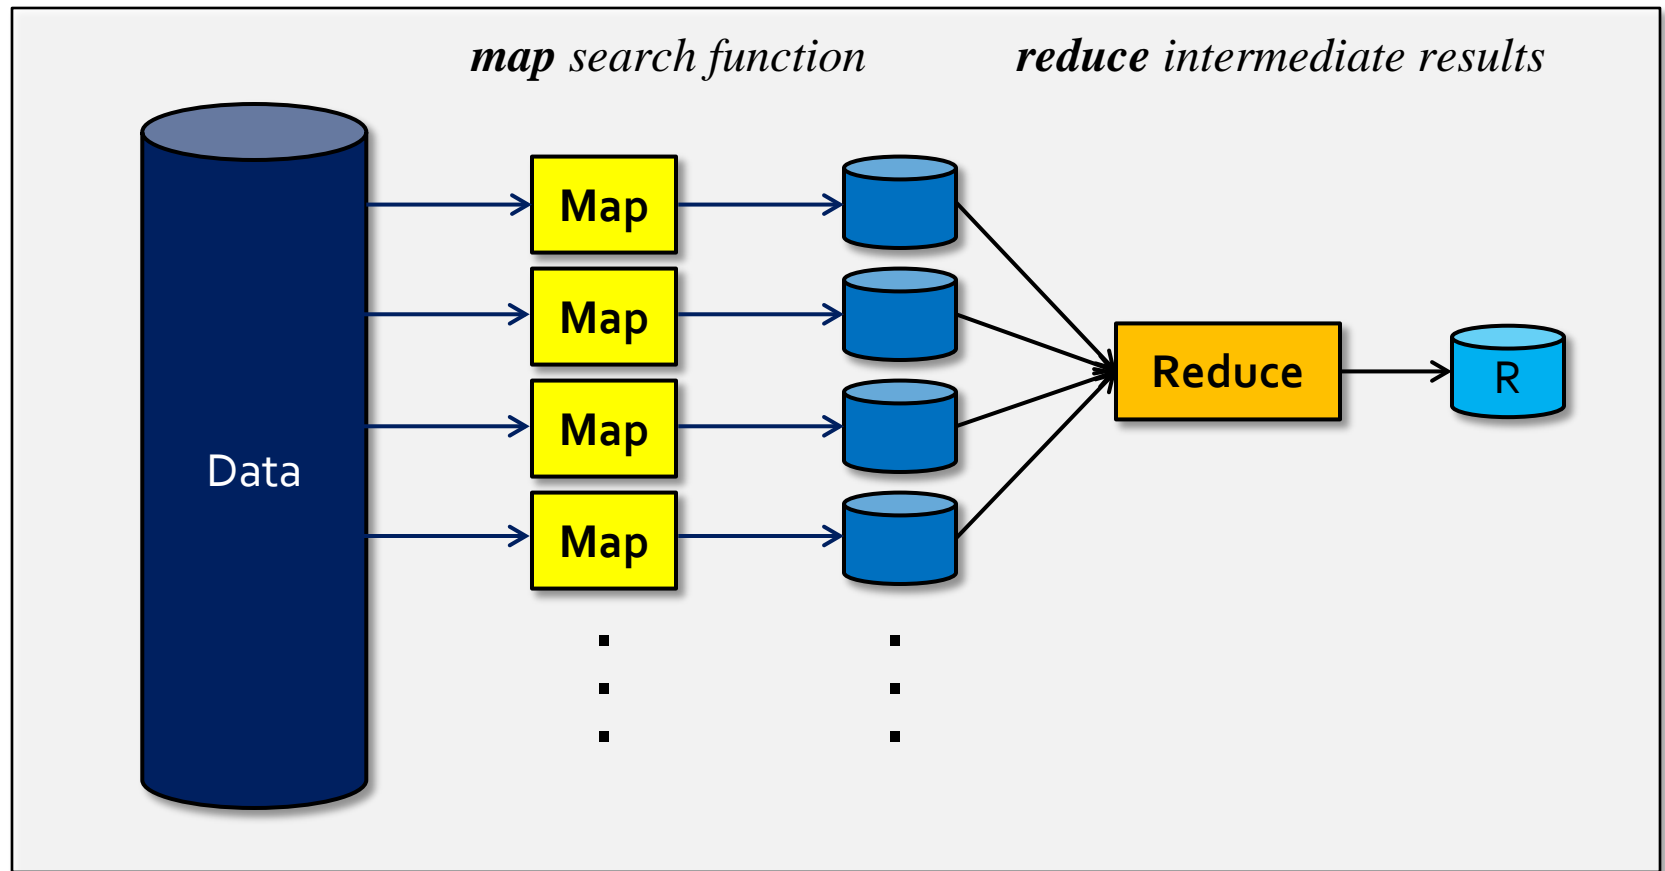
- However, improvement was **sub-linear**

- *Bad fit given that I/O time and processing times are roughly the same (fast)?*
- *Contention for the blocking collection?*

Let's find out by applying another pattern  
— MapReduce — that eliminates  
contention by eliminating shared resources

# MapReduce

- Commonly used for embarrassingly parallel search / data-mining...



# Implementing MapReduce

## ■ Various strategies...

- *fire off  $N$   $\text{Task}<T>$ , use  $\text{WaitAllOneByOne}$  to reduce results as they finish*
- *use  $\text{Parallel.For} / .\text{Foreach}$  with Task Local Storage (TLS)*

```
for (int i=0; i<N; ++i)
    tasks.Add(Task.Factory.StartNew<T>(
        (data) => { return map(data); }
    ));

while (tasks.Count > 0) // wait for tasks to finish:
{
    int index = Task.WaitAny( tasks.ToArray() );
    reduce(tasks[index].Result);
    tasks.RemoveAt(index);
}
```

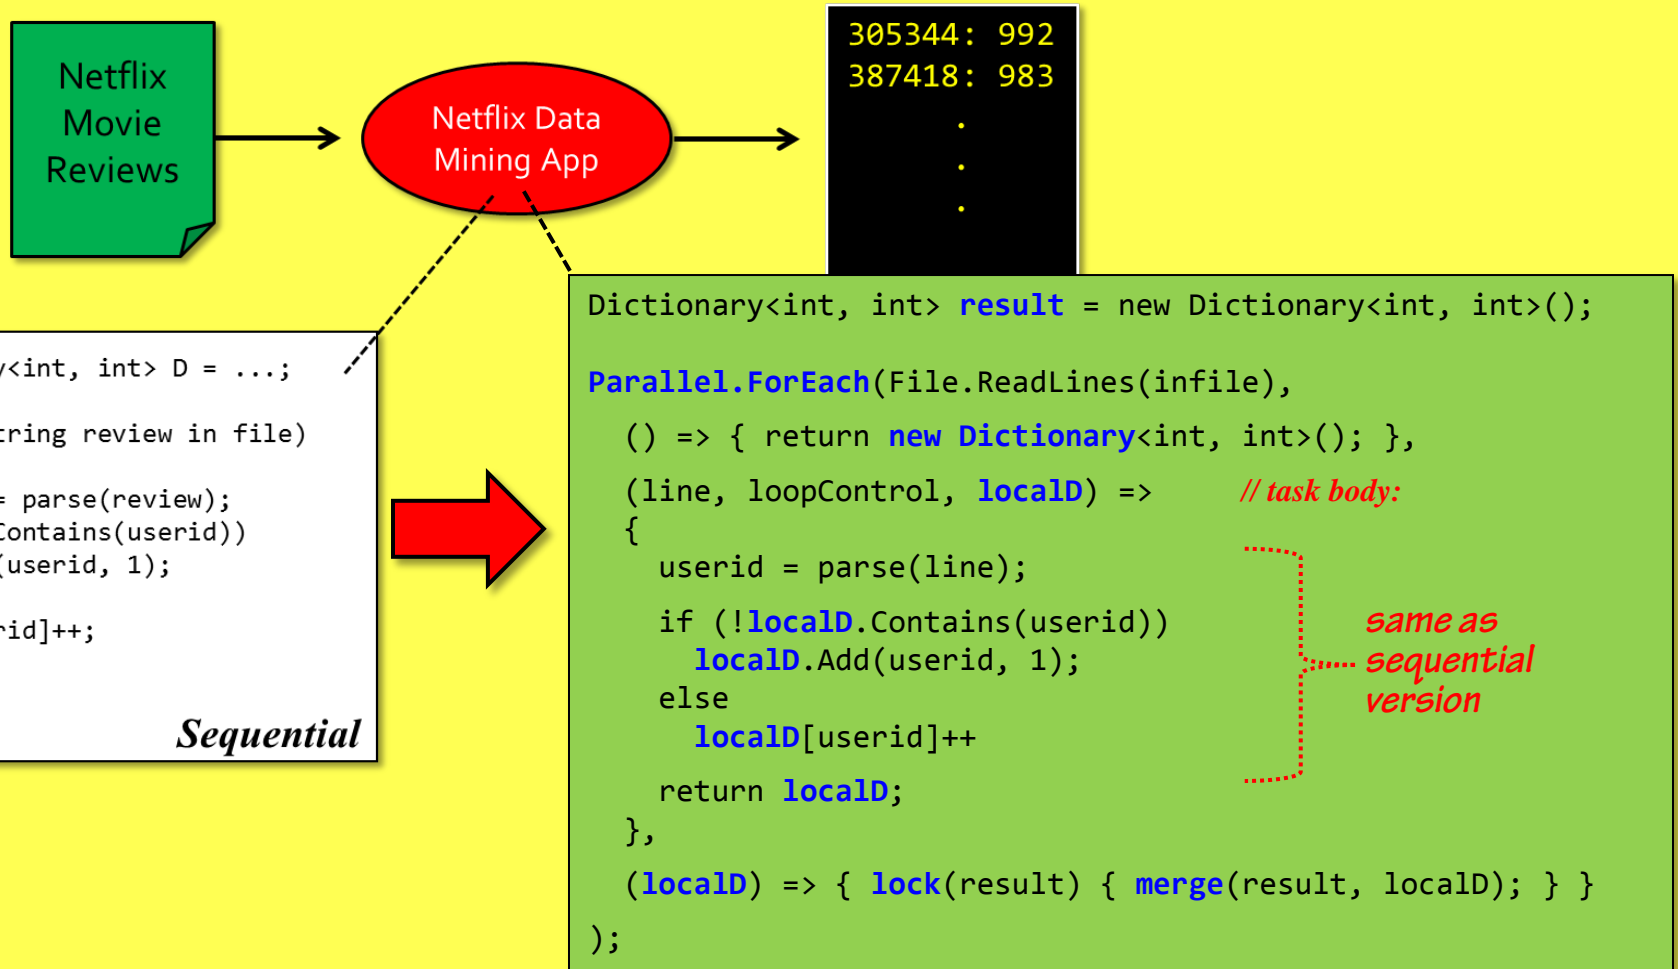
```
Parallel.ForEach(datasource,
    () => { return new TLS(); }, // initializer:

    (datum, ..., tls) => // task body:
    {
        map(datum, tls);
        return tls;
    },

    (tls) => { reduce(tls); } // finalizer:
);
```

# DEMO (v3)

## ■ Netflix data mining using MapReduce...



# Results?

- MapReduce yields **best performance** so far!



- *Eliminating contention usually does...*


## *Extra credit for the viewer!*

*Take a look at the provided solution “Netflix-Par-FileChunk-LessStrict”. It uses MapReduce but improves performance by partitioning the data itself, \*and\* also by a willingness to accept the occasional imprecise result.*



# Parallel LINQ

- Are you a fan of **LINQ** — Language Integrated Query?
- Then you'll love **PLINQ** and its support for:
  - *Parallelism*
  - *MapReduce*



```
var query = ( from x in values.AsParallel()
               where Filter(x)
               select Compute(x) ).Sum();
```

```
var query = ( from x in values
               where Filter(x)
               select Compute(x) ).Sum();
```

# A brief overview...


- LINQ enables querying of **IEnumerable** data sources
- PLINQ turns an **IEnumerable** into a **ParallelQuery** data source
  - *for in-memory collections*

```
var q = (from x in values   where Filter(x) select Compute(x)).Sum();
```

operations that follow `.AsParallel()` run in  
parallel (in most cases)




including `.Sum()`, which performs a parallel (tree-like) reduction





# DEMO (v4)

- Netflix data mining using PLINQ...



```
var ReviewsByUser = File.ReadLines(infile).
    Select(line => parse(line)).
    GroupBy( userid => { return userid; } ).
    Select( g => new { UserId = g.Key, NumReviews = g.Count() } );

var Top10 = (from user in ReviewsByUser
    orderby user.NumReviews descending, user.UserId ascending
    select user).Take(10).ToList();
```

# Speculative Execution


- When you have multiple sources for generating a result...
- **Example:**
  - *Calling web services...*

```
var cts = new CancellationTokenSource();
var token = cts.Token;

for (int i=0; i<N; ++i)
    tasks.Add(Task.Factory.StartNew<T>(
        () => { ... },
        token
    ));

int winner = Task.WaitAny( tasks.ToArray() );
var result = tasks[winner].Result;

cts.Cancel(); // tell the rest to cancel...
```



# Asynchronous Programming Model

- **APM pattern** is commonly used for asynchronous operations
  - *File I/O, network I/O, ...*
  - *“**Begin**” starts operation / “**End**” completes operation & harvests result*
- **Various .NET classes support the APM pattern**
  - **FileStream:**        *BeginRead / EndRead*
  - **HTTPWebRequest:**   *BeginGetResponse / EndGetResponse*
  - ...
- **Advantage?**
  - *.NET classes start operation on a thread, but then return thread to pool until operation completes*

# TPL support for APM pattern

- TPL provides first-class APM support via **façade tasks**
  - Use *FromAsync* to create task that wraps calls to *Begin / End*
  - Use *standard task mechanisms* to wait / continue / harvest result

```
FileInfo fi = new FileInfo(filename);
int bytes  = (int) fi.Length;
byte[] buf = new byte[bytes];

FileStream fs = new FileStream(filename, FileMode.Open, FileAccess.Read,
                               FileShare.Read, bytes, true /*async*/);

var T1 = Task<int>.Factory.FromAsync(fs.BeginRead, fs.EndRead, buf, 0, bytes, null);

var T2 = T1.ContinueWith((antecedent) =>
{
    fs.Close();
    int bytesRead = T1.Result;
    .
    . //process data in buf:
    .
});
```

# DEMO

- Historical **stock data** analysis...

```
C:\Windows\system32\cmd.exe
>> Please enter stock symbol(s) (e.g. 'msft,intc,...'): aapl,intl,msft
** CSV Stock History App [any-cpu, release] **
Stock symbol(s): 'aapl,intl,msft'
Time period: last 10 years
Internet access? True

** aapl **
Data source: 'http://nasdaq.com, daily Close, 10 years'
Data points: 2,515
Min price: $6.56
Max price: $363.13
Avg price: $100.69
Std dev/err: 97.964 / 1.953

** intl **
Data source: 'http://moneycentral.msn.com, weekly Close, 1 year'
Data points: 53
Min price: $15.87
Max price: $26.48
Avg price: $22.21
Std dev/err: 3.452 / 0.474

** msft **
Data source: 'http://moneycentral.msn.com, weekly Close, 1 year'
Data points: 53
Min price: $23.27
Max price: $28.60
Avg price: $25.73
Std dev/err: 1.362 / 0.187

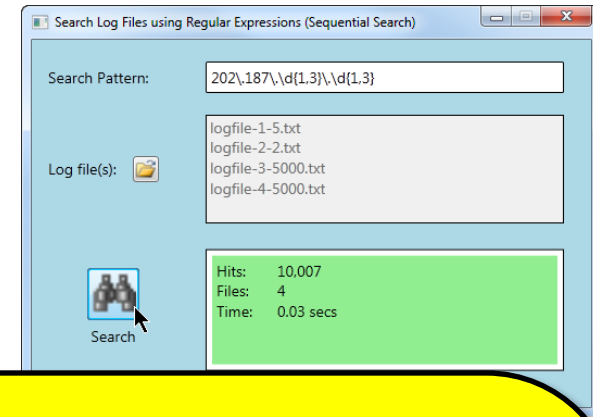
** Done **
Press any key to continue . . .
```

*Good example of both patterns:*

- 1. Speculative execution of downloads*
- 2. APM and façade tasks for HTTP requests*

# Aside: Parallel I/O


- I/O is **\*hard\*** to parallelize
  - *requires support from underlying OS and physical hardware*
  - *sensitive to data sizes*
  - *don't expect much performance from commodity hardware*



## *Design Challenge!*

*The first lecture — “Understanding the Dangers of Concurrency” — focused on a Log File Search app. However, none of the provided solutions work for arbitrarily large files (they all fail with an “Out of Memory” exception).*

*Provided in the “before\” sub-folder of this lecture is a sequential version that works for arbitrarily large files, along with two parallel versions that fail (one does synchronous I/O, and the other does async I/O using the APM). Based on what you’ve learned in this course, design a correct, high-performing parallel version for arbitrarily large files.*



# Summary

- Software design is **challenging**
- Asynchronous and Parallel software is **even more challenging :-)**
- Success measured in terms of both **correctness** and **performance**
- Increase your odds of success by using patterns:
  - *Structured parallelism*
  - *Pipeline and Dataflow*
  - *Concurrent data structures*
  - *Producer-Consumer*
  - *MapReduce*
  - *PLINQ*
  - *Speculative execution*
  - *APM*

# References

- **Microsoft's main site for all things parallel:**
  - <http://msdn.microsoft.com/concurrency>
- **LINQ:**
  - LINQ: <http://code.msdn.microsoft.com/101-LINQ-Samples-3fb9811b>
  - PLINQ: <http://msdn.microsoft.com/en-us/library/dd460688.aspx>
  - Parallel.Foreach vs. PLINQ:  
<http://blogs.msdn.com/b/pfxteam/archive/2010/04/21/9997559.aspx>
- **I highly recommend the following short, easy-to-read book:**
  - *Parallel Programming with Microsoft .NET: Design Patterns for Decomposition and Coordination on Multicore Architectures*, by C. Campbell, R. Johnson, A. Miller and S. Toub, Microsoft Press

**Online:** <http://tinyurl.com/tpl-book>